

Building a Scalable 3-Tier Data Lakehouse for Mobility Analysis in Spain

Big Data Engineering Project

Joan Sánchez Verdú María López Hernández
Fernando Blanco Membrives

Link to Code Repository:

<https://github.com/joanstudyai-code/MUCEIM-BDET-Project.git>

December 27, 2025

Abstract

Abstract. This paper presents the design and implementation of a scalable 3-tier Data Lakehouse architecture tailored for the analysis of large-scale mobility data in Spain. Leveraging open data from the Spanish Ministry of Transport (MITMA) and the National Statistics Institute (INE), the system integrates high-volume Origin-Destination (OD) matrices with demographic and economic indicators. The infrastructure utilizes DuckDB for efficient analytical processing and DuckLake for ACID-compliant storage, orchestrated via Apache Airflow. We demonstrate the system's utility through two key use cases: characterizing typical daily mobility patterns and identifying transport infrastructure gaps using gravity models. The result is a robust, cost-effective platform enabling transport experts to derive data-driven insights for urban planning.

1 Introduction

Mobility analysis is critical for modern urban planning...

1.1 Objectives

The primary objective of this work was to construct a robust data infrastructure...

- Build a scalable Lakehouse for mobility and economic data.
- Implement ELT processes across Bronze, Silver, and Gold tiers.
- Demonstrate analytical value through specific mobility use cases.

2 Data Sources

The data lakehouse architecture is designed to ingest and integrate public domain data exclusively, ensuring reproducibility and open access. The following subsections detail the specific datasets and sources selected to build the information system.

2.1 MITMA Open Data

The core mobility dataset is sourced from the Spanish Ministry of Transport, Mobility and Urban Agenda (MITMA [1]), specifically the *Estudios Básicos de Movilidad*. This dataset provides high-resolution insights into the daily movements of residents derived from mobile network big data. From the MITMA the following files have been used:

- **(202301-202312)_Viajes_municipios.tar:** These files, located in the MITMA web at *estudios_basicos/por-municipios/viajes/meses-completos/*, contain the information of the mobility from an origin to a destination municipality, per day, per hour, for the year 2023.
- **nombres_municipios.csv:** This file, located in the MITMA web at *zonificacion/zonificacion_municipios/*, contains the name of each municipality.
- **relacion_ine_zonificacionMitma.csv:** This file, located in the MITMA web at *zonificacion/*, contains the mapping between the MITMA codes and the INE codes for the municipalities.
- **poblacion.csv:** This file, located in the MITMA web at *zonificacion/*, contains the population of each municipality per year.
- **zonificacion_municipios.shp:** This ESRI Shapefile (along with its *.dbf*, *.shx*, and *.prj* components), located in the MITMA web at *zonificacion/zonificacion_municipios/*, contains the official geospatial geometries (polygons) for the municipal study zones.

2.2 INE Demographics and Economics

To contextualize mobility flows and enable the gravity model analysis, mobility data is enriched with official statistics from the Spanish National Statistics Institute (INE [2]). The statistics used from INE are the following:

- **ine_rent_municipalities.csv:** This file, located in the INE web at https://www.ine.es/jaxiT3/files/t/es/csv_bd/, contains the rent information per year and municipality.

2.3 Open Data

To categorize mobility patterns by day type (e.g., distinguishing standard workdays from holidays), specific calendar data is integrated from Open Data portals.

- **calendario_laboral.csv:** This file, sourced from the Madrid City Council Open Data Portal (*Datos Abiertos Madrid* [4]) at <https://datos.madrid.es/>, contains the official working calendar. It provides the classification of dates (working days, weekends, and holidays) from 2013 to 2026 (included).

3 Lakehouse Architecture

The system follows the Medallion Architecture pattern, decoupled into storage and compute layers to ensure scalability.

3.1 Technology Stack

The platform is built upon a modular, open-source stack designed to decouple compute from storage, allowing for independent scaling and cost efficiency. The architecture leverages the following core components:

- **Compute Engine (DuckDB [5]):** Selected for its high-performance, vectorized execution engine optimized for OLAP workloads. Unlike distributed systems (e.g., Apache Spark), DuckDB runs in-process, significantly reducing infrastructure overhead while maintaining the ability to process datasets larger than RAM efficiently.
- **Lakehouse Management (DuckLake [6]):** A specialized extension that brings Data Lakehouse capabilities to the ecosystem. It provides ACID transaction support (Atomicity,

Consistency, Isolation, Durability) and snapshot isolation. This ensures that long-running ingestion tasks do not block analytical queries and prevents data corruption during partial write failures.

- **Storage - Cloud Infrastructure (Neon + S3):** The system adopts a cloud-native architecture that physically separates data from metadata:
 - **Data Plane (AWS S3):** Stores the actual physical files (Parquet and CSV). S3 provides durability and scalable object storage for the Bronze, Silver, and Gold layers.
 - **Control Plane (Neon Postgres):** A serverless PostgreSQL instance acts as the centralized catalog and metadata store. It manages table definitions, schemas, and transaction locking, enabling safe concurrent writes to S3 from multiple Airflow workers.
- **Orchestration (Apache Airflow [7]):** Manages the end-to-end lifecycle of data pipelines. Airflow allows for the definition of Directed Acyclic Graphs (DAGs) to handle complex task dependencies, backfilling strategies for historical data, and granular retries in case of network or source failures.

3.2 Data Layering Strategy

The data pipeline is structured following the standard “Medallion” architecture pattern (Multi-hop), which organizes data quality into progressive stages. This design ensures that raw data is preserved for auditability while downstream layers provide cleaned and enriched datasets optimized for analytical consumption. The architecture consists of three distinct zones, each with a specific purpose and schema design.

3.2.1 Bronze Layer (Raw)

The Bronze layer acts as the landing zone (Staging Area) for all external data. Its primary function is to capture data from source systems in its native format with minimal modification, ensuring an immutable record of the original input.

In this implementation (Figure 1), a distinct table is instantiated for each source entity to maintain data isolation. To ensure pipeline robustness and prevent ingestion failures due to type mismatches, tabular data columns (such as CSV files) are initially cast as `VARCHAR`. A notable exception is the geographic data: by utilizing DuckDB’s `spatial` extension, the ESRI Shapefile components are ingested via the `ST_Read` function. This allows the system to capture municipal boundaries directly as native `GEOMETRY` objects, preserving spatial precision from the source.

Additionally, the ingestion process is configured to ignore malformed records, ensuring that isolated errors do not halt the entire pipeline. To enhance data governance and auditability, metadata columns are appended to every record, including `ingestion_timestamp` and `source_url`. Finally, to optimize storage and performance, the high-volume mobility data is physically partitioned by date (`fecha`), enabling efficient batch processing.

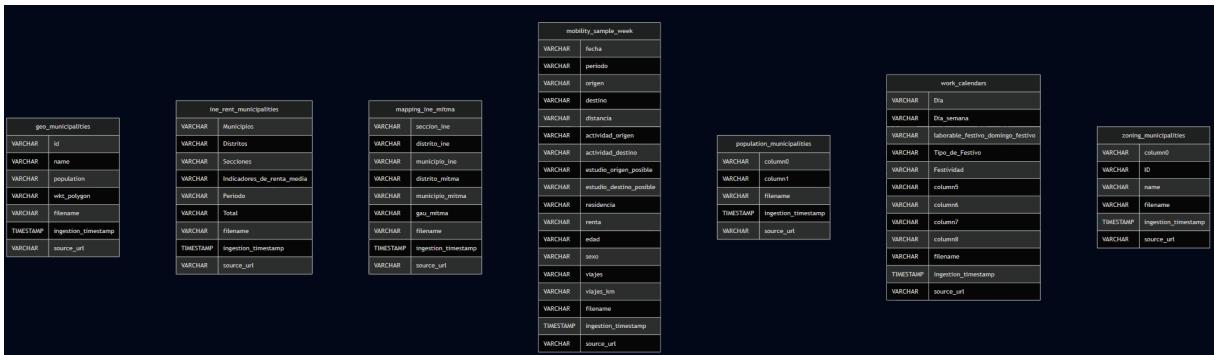


Figure 1: Entity Diagram of the Bronze Layer. One table for each file.

3.2.2 Silver Layer (Cleaned & Integrated)

The Silver layer represents the refined, enterprise-wide view of the data. In this stage, data undergoes validation, cleaning, standardization, and enrichment processes to transform raw inputs into structured, trustworthy tables. The schema design (Figure 2) follows a Star Schema approach, establishing a central fact table for mobility flows linked to surrounding dimension tables for spatial and socio-economic context. All tables include a **TIMESTAMP WITH TIME ZONE** column **processed_at** for auditability.

Dimension Modeling and Spatial Integration (dim_zones)

The foundation of the Silver layer is the creation of a unified spatial dimension, `dim_zones`. This table integrates the descriptive zoning metadata (MITMA codes, INE codes, and municipality names) with the geospatial vector data ingested in the Bronze layer.

- **Spatial Integration:** The `dim_zones` table leverages the native `GEOMETRY` objects already parsed in the Bronze layer. By joining the municipal geometries with the MITMA-INE crosswalk (`mapping_ine_mitma`), the Silver layer establishes a clean, spatially-enabled dimension.
 - **Surrogate Keys:** A unique integer identifier (`zone_id`) is generated for each municipality. This decouples downstream analytics from changes in external string codes (MITMA/INE IDs) and improves join performance.

Socio-Economic Metrics and Filtering

Complementary metric tables are generated by cleaning raw CSVs. For the `metric_population`, data is cast to `BIGINT` while removing metadata headers.

For the `metric_ine_rent` table, specific business logic is applied to the INE source file to resolve granularity mismatches. The pipeline filters the dataset to isolate the “Average Net Income per Person” indicator and explicitly excludes district-level or census-section entries (filtering out rows where `Distritos` or `Secciones` are populated). This ensures that the resulting economic metrics align perfectly with the municipal granularity of the `dim_zones` table.

Temporal Context (Holidays)

To support temporal analysis, a `dim_zone_holidays` table is constructed. This process involves filtering the raw working calendar for events classified as “National Holidays” and cross-referencing them with the zones. This dimension allows the Gold layer to accurately distinguish mobility patterns between standard working days and holidays.

Mobility Fact Table Construction

The core mobility data is transformed into the `fact_mobility` table. This process converts the raw, text-based daily CSVs into a strongly typed, time-series optimized format. Key transformations include:

1. **Temporal Standardization:** The raw separate date and hour fields are combined into a single `TIMESTAMP WITH TIME ZONE` column, explicitly localized to 'Europe/Madrid'.
2. **Spatial Lookup:** Origin and Destination codes are replaced by their corresponding `zone_id` through double joins against the `dim_zones` table, ensuring referential integrity.
3. **Type Casting:** The "trips" column is cleaned (handling European decimal formats) and cast to `DOUBLE` precision to support fractional trip expansion factors.

Data Observability and Quality Logging

To ensure the reliability of the pipeline, a dedicated table named `data_quality_logs` has been implemented within the Silver layer. Unlike the business data tables, this table serves as a metadata registry to persist the results of automated audit checks performed during ingestion (e.g., null rate calculations, row count validations, or statistical anomalies). This allows for longitudinal tracking of data health.

The schema is designed to be generic enough to store heterogeneous metrics:

- `check_timestamp`: The exact time when the audit was executed.
- `table_name`: The target table being audited (e.g., `fact_mobility`).
- `metric_name`: A descriptive identifier for the KPI (e.g., `avg_income_per_capita`, `null_zone_rate`).
- `metric_value`: A numeric field (`DOUBLE`) storing the result of the check.
- `notes`: Textual field for context, such as the specific batch date range or error warnings associated with the metric.

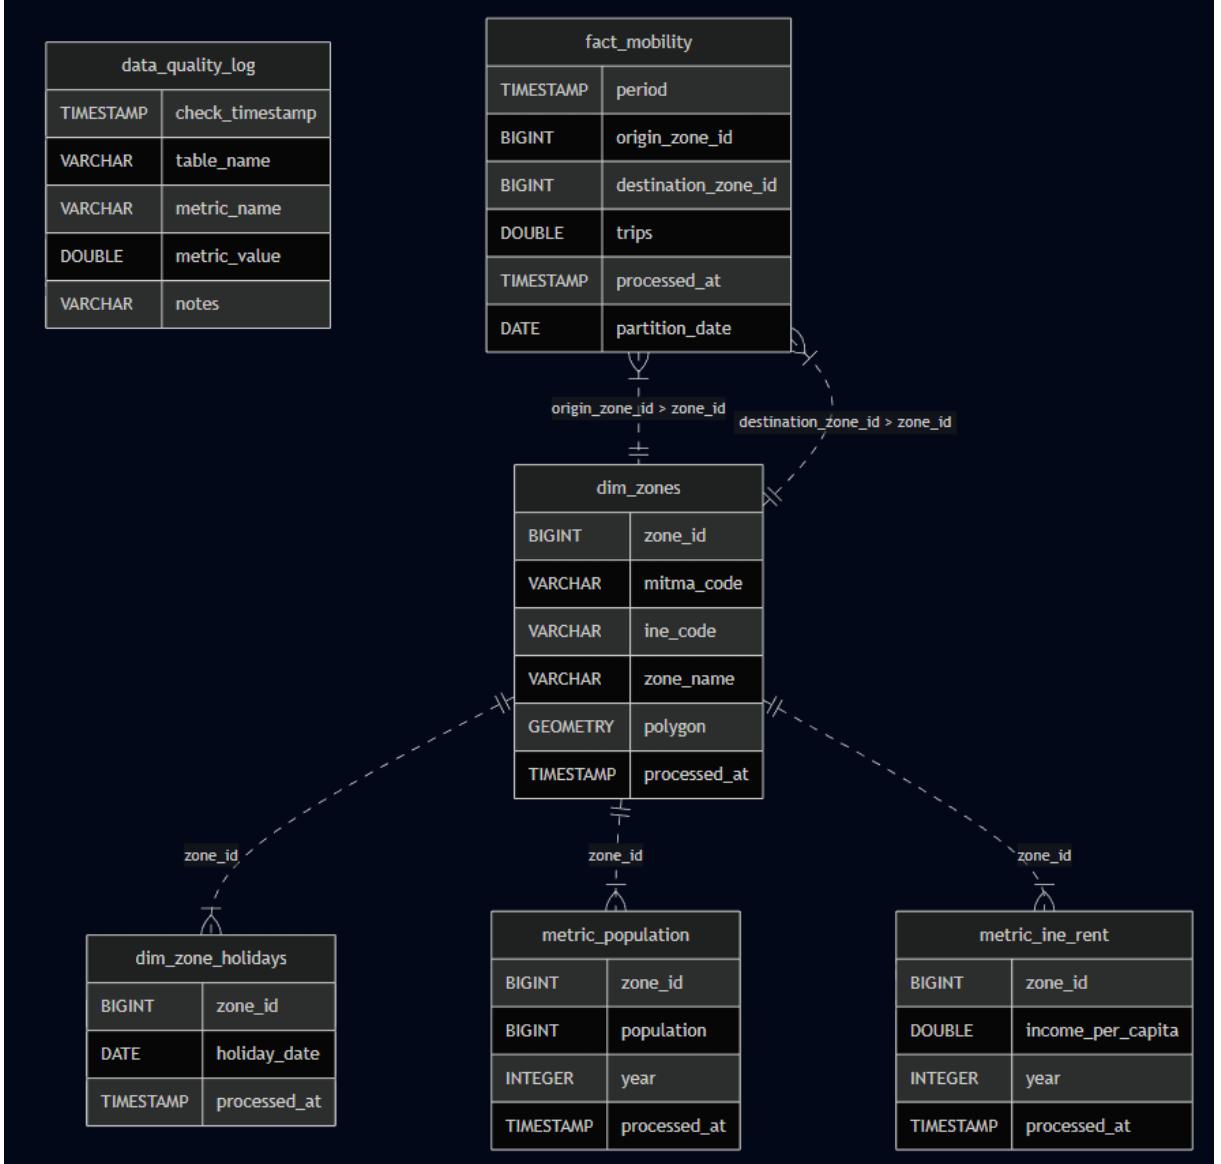


Figure 2: Entity-Relationship Diagram (ERD) of the Silver Layer. The schema follows a Star Schema design, centering on the `fact_mobility` table linked to spatial (`dim_zones`) and temporal dimensions.

3.2.3 Gold Layer (Mart)

[TODO]

4 Implementation Methodology

The operational logic of the Data Lakehouse is implemented through a modular, automated workflow orchestrated by Apache Airflow [7]. The system adopts a decoupled architecture where distinct Directed Acyclic Graphs (DAGs) handle specific stages of the data lifecycle based on their update frequency and functional scope.

This section details the engineering strategies employed to ensure scalable data ingestion, atomic processing of daily batches, and the automated generation of analytical models. The implementation prioritizes robustness through strict transaction management, task isolation, and the separation of heavy computational ELT from on-demand reporting queries.

4.1 Orchestration Strategy

To manage the dependencies between the Bronze, Silver, and Gold layers efficiently, the system is architected into a multi-DAG ecosystem. This design allows for independent scaling and maintenance of static dimensions versus high-volume mobility facts.

The workflow is organized into a hierarchical structure controlled by a Master Orchestrator, as described below:

1. **Infrastructure & Dimensions DAG:** Handles the ingestion of low-frequency static data (INE demographics, MITMA zoning, Calendars). It establishes the schema foundations and Bronze/Silver dimension tables.
2. **Mobility Ingestion DAG:** A parameterized worker pipeline designed for high-volume processing. It accepts date ranges to ingest, clean, and transform the daily mobility files (MITMA OD Matrices) from Bronze to Silver using atomic tasks.
3. **Gold Generation DAG:** Triggered upon the completion of ingestion, this pipeline aggregates the Silver data to construct the core analytical models (e.g., K-Means Clustering, Gravity Models) in the Gold layer.
4. **Business Reporting DAGs:** A set of independent, on-demand DAGs designed to query the pre-computed Gold layer. These are isolated from the ELT process to allow “Transport Experts” to generate specific reports (e.g., specific Polygon filters or time windows) without re-processing the underlying data.
5. **Master Orchestrator DAG:** A controller pipeline containing no logic other than the triggering of downstream DAGs. It manages the sequential dependencies between infrastructure setup, data ingestion, and model training, ensuring referential integrity across the Lakehouse.

4.2 Infrastructure & Dimensions (DAG 1)

The foundational pipeline, `infrastructure_and_dimensions`, is responsible for establishing the Lakehouse schema and ingesting low-velocity reference data. Unlike the daily mobility processing, this DAG is designed to be executed on-demand or annually, as zoning definitions and census statistics change infrequently.

The workflow is architected into three distinct logical phases to maximize parallelism while ensuring referential integrity.

4.2.1 Phase 1: Infrastructure Initialization

Before data movement begins, the pipeline ensures the environment is consistent. The task called `create_schemas` connects to the Neon catalog to verify the existence of the logical namespaces (`bronze`, `silver`, `gold`). Simultaneously, the `create_stats_table` task initializes the `data_quality_log` table in the Silver layer, which is a prerequisite for audit logging in all subsequent pipelines.

4.2.2 Phase 2: Parallel Bronze Ingestion

To optimize total execution time, all external data extraction tasks are configured to run in parallel. These tasks use the Airflow `@task` notation arguments to specify 3 retries with a gap of 1 minute in case they fail. This phase handles heterogenous data formats through specialized ingestion logic:

- **Spatial Ingestion:** The `br_ingest_geo_data` task utilizes DuckDB's `spatial` extension. It performs an HTTP HEAD request to validate the existence of the remote Shapefiles (.shp, .shx, .dbf) before executing `ST_Read` to ingest the vector geometry directly into the Bronze layer.
- **Statistical Ingestion:** Tasks targeting INE data (e.g., `br_ingest_ine_rent`) implement specific handling for legacy formats, forcing ISO-8859-1 encoding and tab separators to prevent parsing errors common in Spanish government datasets.
- **Dictionaries:** Tasks for zoning and mapping ingest standard CSVs via DuckDB's `read_csv_auto`, adding audit columns (`source_url`, `ingestion_timestamp`) on the fly.

4.2.3 Phase 3: Silver Transformation and Convergence

The dependency structure in the Silver layer follows a "Funnel" pattern (Figure 3). The system enforces a hard dependency on the creation of the master dimension table before populating satellite metric tables.

1. **Convergence Point (`dim_zones`):** This task acts as the synchronization barrier. It waits for the completion of three specific Bronze tasks—Geometry, Zoning names, and INE Mapping. Once available, it performs a 3-way join to construct the master `dim_zones` table, generating the surrogate `zone_id`.
2. **Satellite Expansion:** Once the master dimension exists, the downstream tasks called `metric_population`, `metric_ine_rent` and `dim_zone_holidays` trigger in parallel. Each task performs a join between its respective raw Bronze source and the newly created `dim_zones` to assign the correct surrogate keys to the statistical data.

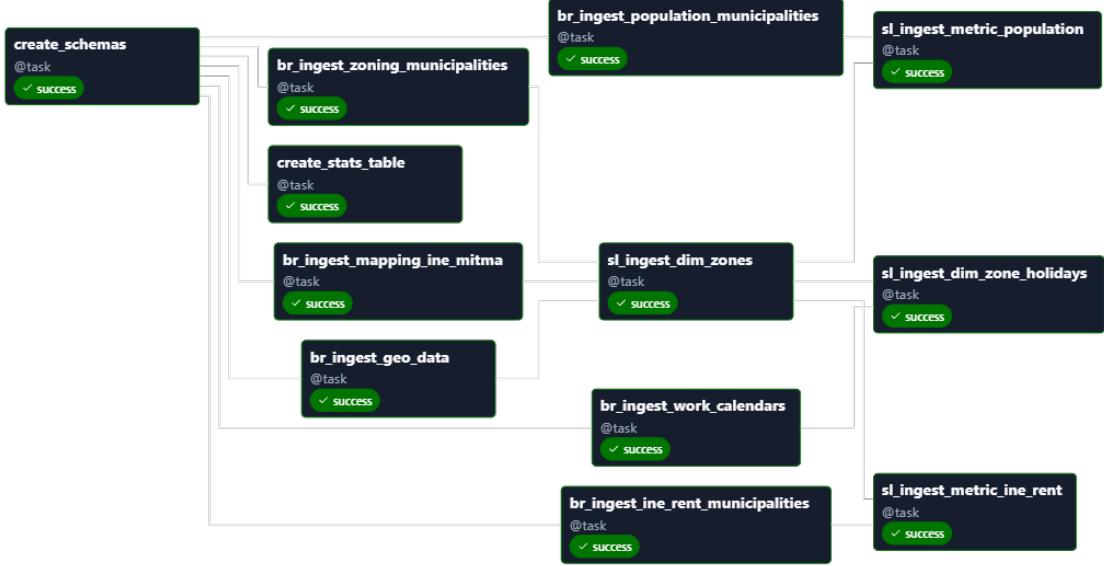


Figure 3: Dependency graph of the Infrastructure DAG. Note the convergence of spatial and dictionary data into `dim_zones` before metric processing.

4.3 Mobility Ingestion (DAG 2)

The second pipeline, `mobility_ingestion`, constitutes the core factory of the Lakehouse. It is responsible for the daily ingestion and transformation of the high-volume Origin-Destination matrices. Given that the source data is partitioned by day (one compressed CSV per day), this DAG utilizes Airflow’s Dynamic Task Mapping to scale horizontally, processing date ranges supplied via run-time parameters.

4.3.1 Phase 1: Dynamic Scheduling and Initialization

Unlike the static infrastructure DAG, this workflow begins by generating a execution plan based on user input. The `generate_date_list` task parses the `start_date` and `end_date` parameters to produce a list of strings (e.g., `['20230101', '20230102', ...]`).

Simultaneously, the pipeline executes idempotent initialization tasks (`ensure_br_mobility...` and `ensure_sl_mobility...`). These tasks utilize `CREATE TABLE IF NOT EXISTS` logic to prepare the partitioning structures in the Bronze and Silver layers, ensuring that parallel workers do not encounter race conditions when attempting to write to non-existent tables.

4.3.2 Phase 2: Resilient Bronze Ingestion

The `br_process_single_day` task is mapped over the generated date list. Each instance handles a specific day, implementing robust checks to handle the known inconsistencies in the MITMA source server:

- **Pre-flight Validation:** Before attempting extraction, the task performs an HTTP HEAD request to the calculated URL. If the source file is missing (a common occurrence for specific dates in late 2023), the task logs a warning and skips execution gracefully rather than failing the pipeline.
- **Throttling and Retries:** To prevent overwhelming the source server or saturating local memory, concurrency is strictly limited via `max_active_tis_per_dag=2`. Network resilience is enforced with a retry policy of 5 attempts with a 30-second delay.
- **Streaming Ingestion:** Valid files are streamed directly from the remote server into the `mobility_data` table using DuckDB's `read_csv_auto`, bypassing local disk storage.

4.3.3 Phase 3: Silver Transformation

The transformation phase (`sl_process_single_day`) mirrors the mapping of the ingestion phase but introduces a strict cross-layer dependency: Silver processing for the batch only begins once Bronze ingestion is confirmed.

Inside this task, raw text data is transformed into the analytical schema. This involves joining the raw mobility records with the `dim_zones` table (created in DAG 1) to resolve MITMA/INE codes into internal surrogate keys. This effectively acts as a foreign key validation constraint; records with invalid zone codes are implicitly filtered during the inner join, ensuring referential integrity in the Silver layer.

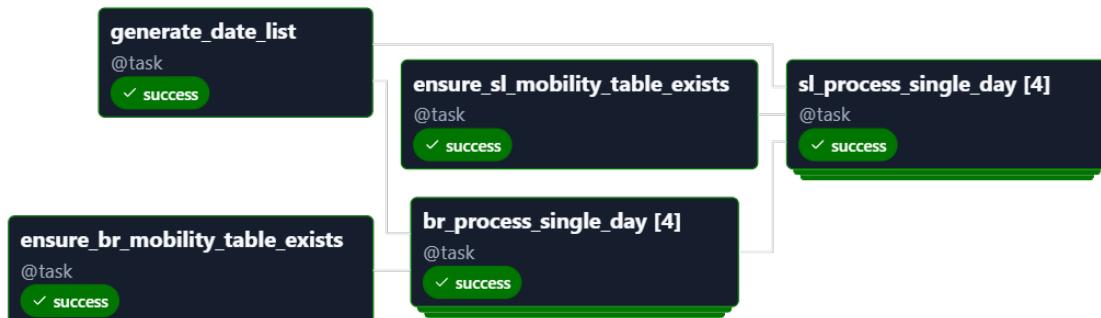


Figure 4: The Mobility Ingestion DAG utilizing Dynamic Task Mapping. The central block expands into N parallel tasks corresponding to the number of days in the requested date range.

4.4 Gold Generation (DAG 3)

[TODO]

4.5 Bussines Reporting (DAG 4-5-6)

[TODO]

4.6 Master Orchestrator (DAG 7)

[TODO]

5 Results and Use Cases

5.1 Use Case 1: Typical Mobility Patterns

To characterize the “typical day”, we aggregated hourly flows and applied K-Means clustering...

Figure 5: Hourly mobility profiles identified (Weekdays vs. Weekends).

5.2 Use Case 2: Infrastructure Gap Analysis

We compared actual MITMA demand against theoretical demand derived from a gravity model:

$$T_{ij} = k \cdot \frac{P_i \cdot E_j}{d_{ij}^2} \quad (1)$$

The mismatch ratio calculated in the Gold layer highlights zones where actual transport usage significantly lags behind potential demand...

6 Discussion

The implemented lakehouse allows experts to query billions of rows using standard SQL without managing infrastructure...

7 Conclusions and Limitations

The project successfully established a scalable 3-tier architecture. However, limitations exist regarding data anomalies in late 2023...

References

- [1] Spanish Ministry of Transport (MITMA). Open Data Mobility. <https://www.transportes.gob.es/ministerio/proyectos-singulares/estudios-de-movilidad-con-big-data/opendata-movilidad>
- [2] Spanish National Statistics Institute (INE). <https://www.ine.es/>
- [3] Centro Nacional de Información Geográfica (CNIG). <https://www.ign.es/>
- [4] Ayuntamiento de Madrid. Portal de Datos Abiertos. <https://datos.madrid.es/>
- [5] DuckDB Documentation. <https://duckdb.org/>

- [6] DuckLake. *DuckLake Documentation*. <https://ducklake.select/docs/stable/>
- [7] Apache Software Foundation. *Apache Airflow Documentation*. <https://airflow.apache.org/>